

AUTO-VECTORIZATION USING AVX

Advanced Programming Tutorial

Date: 10./13.12.2013

Tutor: Felix Rempe

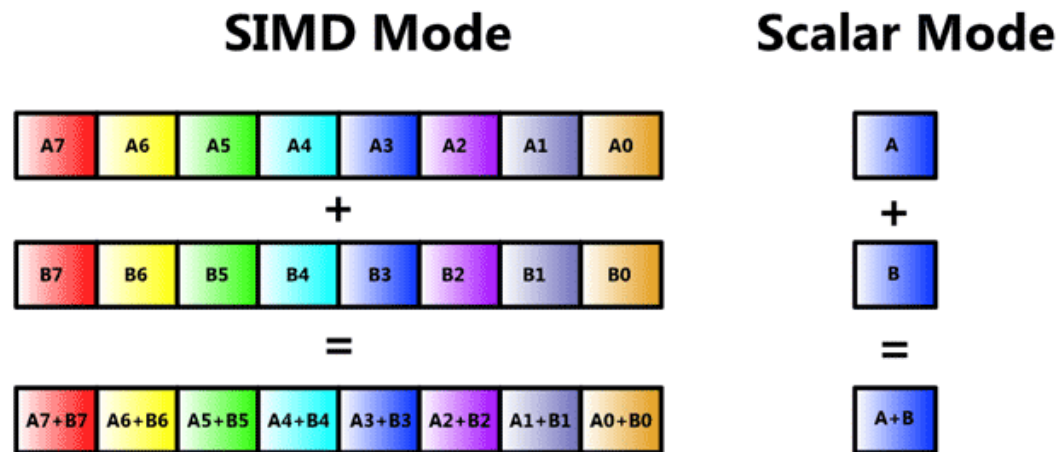
AVX and Auto-Vectorization

- Introduction / Motivation
- How to use
- Auto-Vectorization Requirements
- Useful #pragmas for Vectorization
- Guided Auto-Vectorization
- Summary

INTRODUCTION

Introduction to AVX


- What is AVX? (Advanced Vector Extensions)
 - Set of instructions to do SIMD commands on Intel architectures (beginning with Sandy Bridge)
 - Previous: SSE (Streaming SIMD Extension)
 - Principle: 1 command for vector operations



Introduction to AVX


- Features:

- 256bit registers for operands (only FP)



8x float
4x double

- 128bit for all other types



4x float
2x double
16x byte
word
doubleword
quadword
doublequadword

Source: <http://software.intel.com/>

- Many types of operations available
- Easily computing in e.g. $A = B + C$ (vector + vector)


Motivation: Example Mandelbrot

- Mandelbrot definition:

- C: complex
- N: iterations

$$z_{n+1} = z_n^2 + c$$


All c for which $\lim_{n \rightarrow \infty} z_n$ is bounded belong to the Mandelbrot set



```
1 z,p are complex numbers
2 for each point p on the complex plane
3   z = 0
4   for count = 0 to max_iterations
5     if abs(z) > 2.0
6       break
7     z = z*z+p
8   set color at p based on count reached
```

Motivation: Performance on Mandelbrot

- Tested for different image sizes

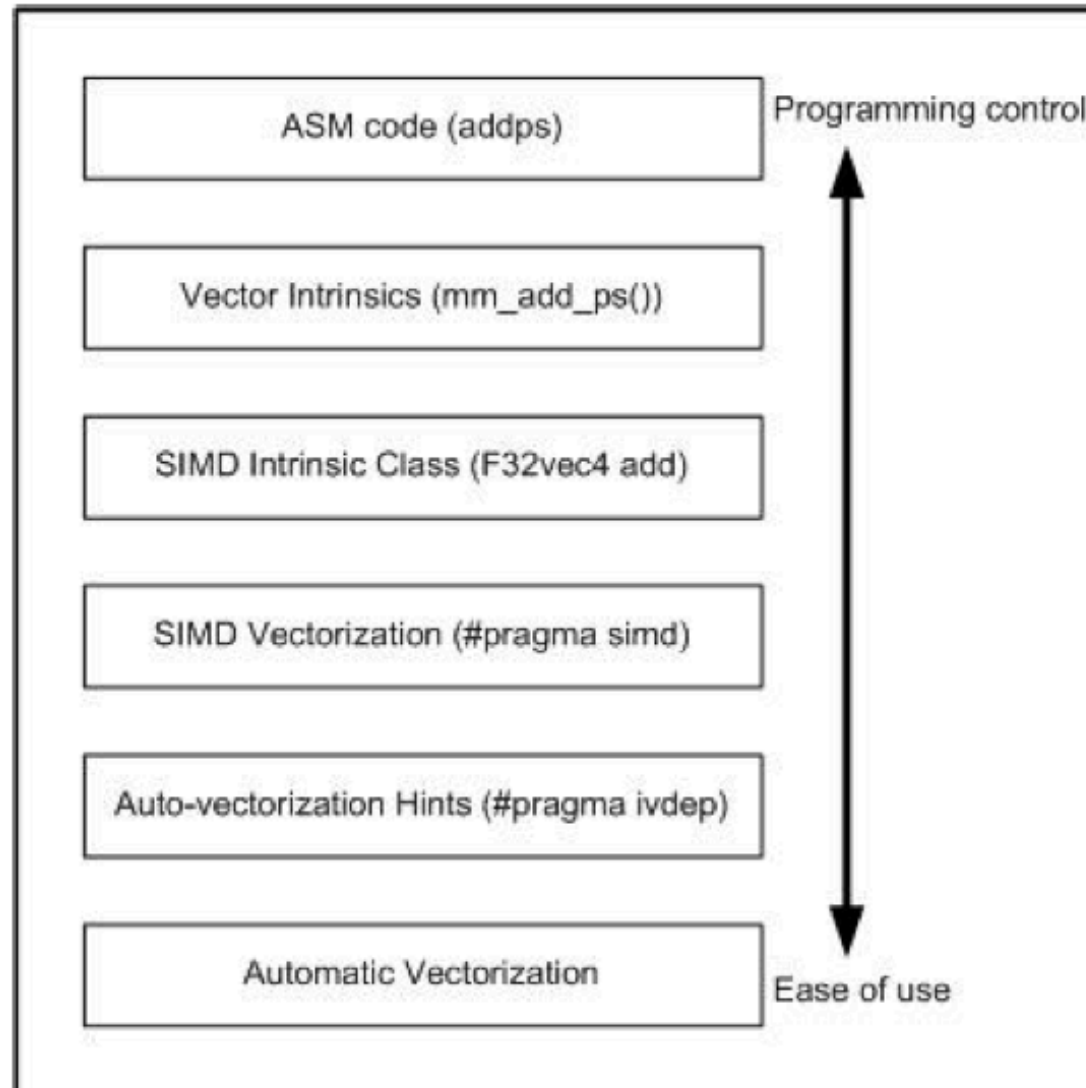


- =>Great speed-up possible!

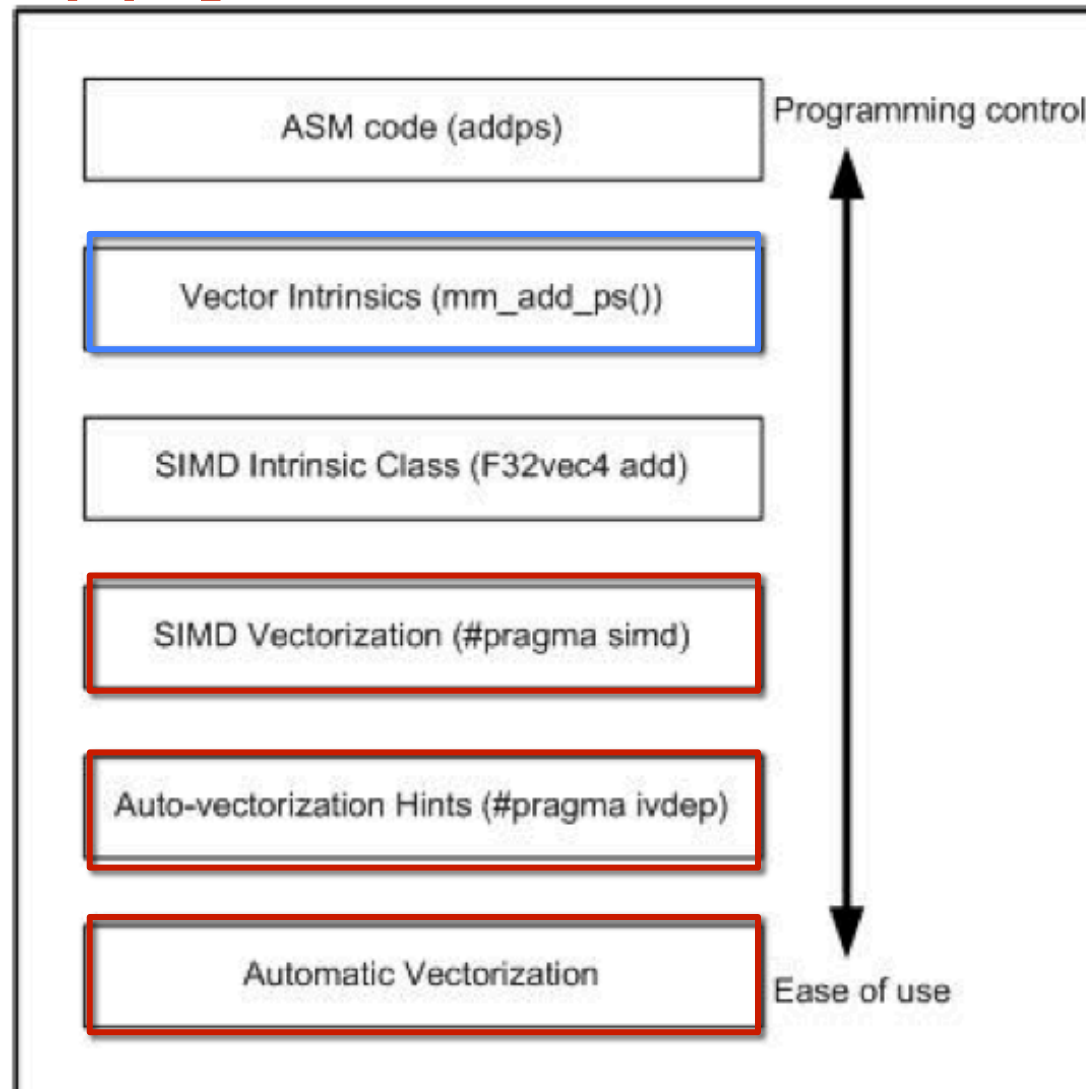
Source:<http://software.intel.com/>

HOW TO USE

How to Apply



How to Apply



How to Apply: Vector Intrinsics

- Example: Load a vector and multiply with another one

```
float incr[8]={0.0f,1.0f,2.0f,3.0f,4.0f,5.0f,6.0f,7.0f};  
__m256 ymm7  = _mm256_load_ps(incr);  
__m256 ymm8  = _mm256_mul_ps(ymm7, ymm0)
```

- In general:

```
_mm256_op_suffix(data_type param1, data_type param2, data_type param3)
```

- Where:
 - *_mm256* is the prefix for working on the 256bit registers
 - *_op* is the operation
 - *_suffix* is the type of data (e.g. s/d for single, double FP)

Auto-Vectorization

- Let the compiler decide what to do
 - No need to fight with the low level commands
 - Portability of code
- Help the compiler by giving *Vectorization Hints*
 - E.g.

```
#pragma vector always  
#pragma ivdep  
...
```

- Automatic Vectorization works best if the user is aware of certain mechanisms

Auto-Vectorization: Compilation

- Commonly used compilers: icc and gcc
- Flags determine how/if vectorization is done
- **icc:**
 - Optimization level: O2 or higher
 - xHost – use the best set of optimization flags for architecture (cf. march= native)
 - xavx – use AVX optimization
- **gcc :**
 - Optimization level: O3
 - march=native – (cf. xHost)
 - mavx – use AVX optimization
 - ftree-vectorize – use vectorization
 - Try to print out the available technologies on your system

```
gcc -march=native -E -v - </dev/null 2>&1 | grep cc1
```

AUTO-VECTORIZATION

Things you should be aware of

Requirements for Auto-Vectorization

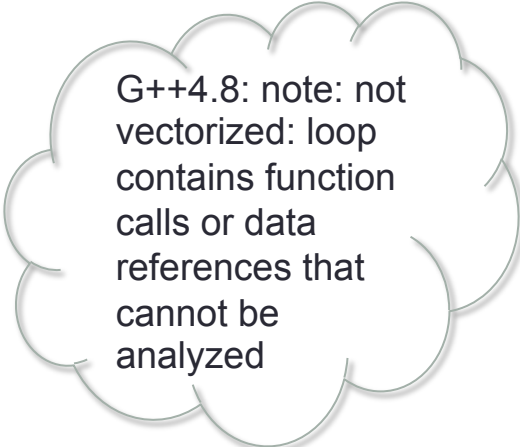
- 1. The number of iterations must be known before entering the loop
- 2. Single entry and single exit
- 3. Conditional expressions might not be vectorized
- 4. Innermost loop
- 5. No function calls (functions that can be inlined are OK)

Example for Rule No. 5

```
float trap_int (float y, float x0, float xn, int nx, float xp, float yp) {  
    float x, h, sumx;  
    for (int i=1;i<nx;i++) {  
        x = x0 + i*h;  
        sumx = sumx + func(x,y,xp,yp);  
    }  
    sumx = sumx * h;  
    return sumx;  
}
```

```
float func(float x, float y, float xp, float yp) {  
    float denom = (x-xp)*(x-xp) + (y-yp)*(y-yp);  
    denom = 1.f/sqrtf(denom);  
    return denom;  
}
```

> icc -c -vec-report2
trap_integ.c
trap_int.c(16) (col. 3):
remark: LOOP WAS
VECTORIZED.



G++4.8: note: not
vectorized: loop
contains function
calls or data
references that
cannot be
analyzed

- Is vectorizable since the function can be inlined and “sqrtf” is vectorizable

More Obstacles

- 1. Non-contiguous memory access:
 - Values that are not adjacent in memory need extra effort to be loaded into registers

```
for (int i=0; i<SIZE; i+=2) {  
    b[i] += a[i] * x[i];  
}
```

```
for (int j=0; j<SIZE; j++) {  
    for (int i=0; i<SIZE; i++) {  
        b[i] += a[i][j] * x[j];  
    }  
}
```

```
for (int i=0; i<SIZE; i+=2) {  
    b[i] += a[i] * x[index[i]];  
}
```

More Obstacles

- 1. Non-contiguous memory access:
 - Values that are not adjacent in memory need extra effort to be loaded into registers

Stride > 1

```
for (int i=0; i<SIZE; i+=2) {  
    b[i] += a[i] * x[i];  
}
```

```
for (int j=0; j<SIZE; j++) {  
    for (int i=0; i<SIZE; i++) {  
        b[i] += a[i][j] * x[j];  
    }  
}
```

```
for (int i=0; i<SIZE; i+=2) {  
    b[i] += a[i] * x[index[i]];  
}
```

More Obstacles

- 1. Non-contiguous memory access:
 - Values that are not adjacent in memory need extra effort to be loaded into registers

Stride > 1

```
for (int i=0; i<SIZE; i+=2) {  
    b[i] += a[i] * x[i];  
}
```

Innermost
loop strides
with j

```
for (int j=0; j<SIZE; j++) {  
    for (int i=0; i<SIZE; i++) {  
        b[i] += a[i][j] * x[j];  
    }  
}
```

```
for (int i=0; i<SIZE; i+=2) {  
    b[i] += a[i] * x[index[i]];  
}
```

More Obstacles

- 1. Non-contiguous memory access:
 - Values that are not adjacent in memory need extra effort to be loaded into registers

Stride > 1

```
for (int i=0; i<SIZE; i+=2) {  
    b[i] += a[i] * x[i];  
}
```

Innermost
loop strides
with j

```
for (int j=0; j<SIZE; j++) {  
    for (int i=0; i<SIZE; i++) {  
        b[i] += a[i][j] * x[j];  
    }  
}
```

Indirect
addressing

```
for (int i=0; i<SIZE; i+=2) {  
    b[i] += a[i] * x[index[i]];  
}
```

More Obstacles

- 2. Data dependencies:
 - If data in the iterations have dependencies, they cannot be safely vectorized

```
for (j=1; j<MAX; j++) {  
    A[j]=A[j-1]+1;  
}
```

```
void test(double * a, double * b) {  
    for (int i = 0; i < SIZE; i++) {  
        a[i] += b[i];  
    }  
}
```

=> Possibility to give the compiler “hints” how to vectorize (using #pragma ivdep)

USEFUL #PRAGMAS

What are Pragmas?

- Pragmas give the compiler additional information what to do
 - E.g. `#pragma once`
- `icc`: lots of pragmas available
 - `#pragma vector always`:
 - Ask the compiler to vectorize even if it is not efficient
 - `#pragma novector`:
 - Don't vectorize
 - `#pragma vector aligned`
 - Asserts that data in the loop is assigned (to 16byte boundaries, for SSE)
 - `#pragma simd`
 - Enforce vectorization
- `gcc(4.8)`: only few

Some Pragmas

- `#pragma ivdep`:
 - Ignore potential data dependencies
 - Assert good data alignment

```
#pragma ivdep
#pragma vector aligned
void test(double * a, double * b) {
    for (int i = 0; i < SIZE; i++) {
        a[i] += b[i];
    }
}
```


Vectorization Pragmas with gcc?

- <http://locklessinc.com/articles/vectorize/>
- Needs some explicit statements:

```
#pragma ivdep
#pragma vector aligned
void test(double * a, double * b) {
    for (int i = 0; i < SIZE; i++) {
        a[i] += b[i];
    }
}
```

icc

```
void test(double * restrict a, double * restrict b) {
    double *x = __builtin_assume_aligned(a, 16);
    double *y = __builtin_assume_aligned(b, 16);
    for (int i = 0; i < SIZE; i++) {
        x[i] += y[i];
    }
}
```

gcc

GUIDED AUTO- VECTORIZATION

Vectorization Reports

- After vectorizing read the compilation report
 - Icc: -vec-report=<n>
 - Gcc: -ftree-vectorizer-verbose=<n>
- Report tells you what has been vectorized
- + what has NOT been vectorized and WHY

n=0: No diagnostic information
n=1: Loops successfully vectorized
n=2: Loops not vectorized - and the reason why not
n=3: Adds dependency Information
n=4: Reports only non-vectorized loops
n=5: Reports only non-vectorized loops and adds dependency info

icc

Level 0: No output at all.
Level 1: Report vectorized loops.
Level 2: Also report unvectorized "well-formed" loops and respective reason.
Level 3: Also report alignment information (for "well-formed" loops).
Level 4: Like level 3 + report for non-well-formed inner-loops.
Level 5: Like level 3 + report for all loops.
Level 6: Print all vectorizer dump information (equivalent to former vect-debug-details).

gcc

Exemplary Report with gcc

- Trying to Auto-vectorize a loop that is not vectorizable

```
g++ -march=native -O3 -ftree-vectorizer-verbose=2 testPProject.cpp -o hallo
```

- GCC4.8 on supermuc returns:

```
float a[N];  
    for ( int i = 0; i < N; i++ ) {  
        a[i] += i;  
    }  
    for ( int i = 1; i < N; i++ ) {  
        a[i] += a[i-1];  
    }
```

Analyzing loop at mandel.cpp:12

mandel.cpp:12: note: not vectorized, possible dependence between data-refs

mandel.cpp:12: note: bad data dependence.

mandel.cpp:12: note: not vectorized, possible dependence between data-refs

mandel.cpp:12: note: bad data dependence.

Guided Vectorization

- Activate with `-guide-vec=<n>` to obtain more information how to vectorize ($n = 1$ or 2)

```
float a[N];  
float c=1;  
for ( int i = 0; i < N; i++) {  
    if ( a[i] > 0 ) {c=a[i]; a[i] = 1 /a[i] ; }  
    if ( a[i] > 1 ) {a[i] +=c;}  
}
```

```
icc -O2 -guide-vec=2 testPProject.cpp -o hallo
```

testPProject.cpp(20): remark #30515: (VECT) Assign a value to the variable(s) "c" at the beginning of the body of the loop in line 20. This will allow the loop to be vectorized. [VERIFY] Make sure that, in the original program, the variable(s) "c" read in any iteration of the loop has been defined earlier in the same iteration.

General hints for Vectorizing

- Use array notation (`[i]`) rather than pointers (`*a + i`)
- Minimize indirect addressing (`a[b[i]]`)
- Try to use loops with unit stride
- Align variables to 16 / 32 byte boundaries

SUMMARY

Summary

- AVX powerful extension to process lots of data
- Automated vectorization is a quick and effective way to get higher performance
- For optimal performance, a deeper understanding of the vectorization basics is necessary
- The GAP and reporting tools of icc help to locate potential

Outlook

- SIMD is becoming more important in the future
 - Performance gain
 - Energy efficiency
- AVX2 has 512bit for vector operations
- Included in the current Haswell architecture